

Google File System

Aman Sharma – 2018201084

Ankit Pant – 2018201035

Ankush Nagpal – 2018201083

Shreya Upadhaya – 2018201091

Outline

Outline

1. Introduction
2. Literature Review
3. Methodology
4. Conclusion
5. Future Scope

Introduction

Introduction

- Google File System – a proprietary file system used by Google

Introduction

- Google File System – a proprietary file system used by Google
- A distributed file system:
 - Hundreds or thousands of inexpensive storage machines

Introduction

- Google File System – a proprietary file system used by Google
- A distributed file system:
 - Hundreds or thousands of inexpensive storage machines
 - Constant monitoring, error detection, fault tolerance, and automatic recovery

Introduction

- Google File System – a proprietary file system used by Google
- A distributed file system:
 - Hundreds or thousands of inexpensive storage machines
 - Constant monitoring, error detection, fault tolerance, and automatic recovery
- Characteristics:
 - Huge files (multi GB) are common

Introduction

- Google File System – a proprietary file system used by Google
- A distributed file system:
 - Hundreds or thousands of inexpensive storage machines
 - Constant monitoring, error detection, fault tolerance, and automatic recovery
- Characteristics:
 - Huge files (multi GB) are common
 - Once written, the files are only read (most often sequentially)

Introduction

- Google File System – a proprietary file system used by Google
- A distributed file system:
 - Hundreds or thousands of inexpensive storage machines
 - Constant monitoring, error detection, fault tolerance, and automatic recovery
- Characteristics:
 - Huge files (multi GB) are common
 - Once written, the files are only read (most often sequentially)
 - File writes are very rare

Introduction

- Google File System – a proprietary file system used by Google
- A distributed file system:
 - Hundreds or thousands of inexpensive storage machines
 - Constant monitoring, error detection, fault tolerance, and automatic recovery
- Characteristics:
 - Huge files (multi GB) are common
 - Once written, the files are only read (most often sequentially)
 - File writes are very rare
 - Most files are mutated by appending new data

Literature Review

- GFS is proprietary and Not for Sale

Literature Review

- GFS is proprietary and Not for Sale
- Hence implementation details are vague

Literature Review

- GFS is proprietary and Not for Sale
- Hence implementation details are vague
- The references mostly outline the architecture and operations

Literature Review

- GFS is proprietary and Not for Sale
- Hence implementation details are vague
- The references mostly outline the architecture and operations
- Recent Changes made to reduce latency not publicly documented yet

GFS Architecture

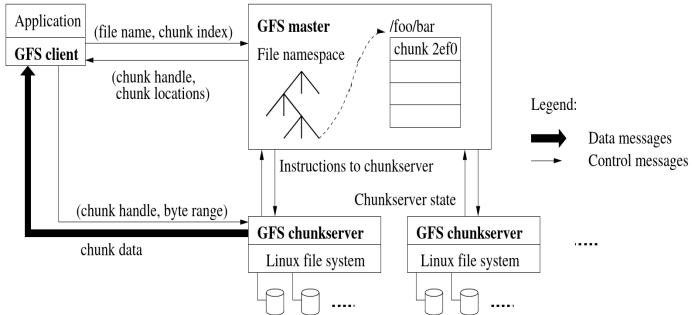


Figure 1: GFS Architecture

The Google File System Architecture [1]

Methodology

- The implementation of GFS was done in python

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File
 - Append

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File
 - Append
 - Snapshot

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File
 - Append
 - Snapshot
 - Load Balancing (chunk replication)

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File
 - Append
 - Snapshot
 - Load Balancing (chunk replication)
 - Garbage Collection

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File
 - Append
 - Snapshot
 - Load Balancing (chunk replication)
 - Garbage Collection
 - Maintaining Integrity

Methodology

- The implementation of GFS was done in python
- The operations supported include:
 - Create File
 - Read
 - Delete File
 - Append
 - Snapshot
 - Load Balancing (chunk replication)
 - Garbage Collection
 - Maintaining Integrity
 - Operation Log

- Create is implemented as:

Create File

- Create is implemented as:
 - Client can issue the create command

Create File

- Create is implemented as:
 - Client can issue the create command
 - Client provides the directory to create file in

Create File

- Create is implemented as:
 - Client can issue the create command
 - Client provides the directory to create file in
 - Client transfers the data to the Master

Create File

- Create is implemented as:
 - Client can issue the create command
 - Client provides the directory to create file in
 - Client transfers the data to the Master
 - Master creates chunks of 64MB from the data (adding padding if required)

Create File

- Create is implemented as:
 - Client can issue the create command
 - Client provides the directory to create file in
 - Client transfers the data to the Master
 - Master creates chunks of 64MB from the data (adding padding if required)
 - Master forwards these chunks to the chunk-servers based on their free disk space

Read

- Read is implemented as:

Read

- Read is implemented as:
 - Client can read chunks of various files using read command

Read

- Read is implemented as:
 - Client can read chunks of various files using read command
 - Client provides the filename to read from

Read

- Read is implemented as:
 - Client can read chunks of various files using read command
 - Client provides the filename to read from
 - Client provides the chunk index (from bytes offset input by user)

Read

- Read is implemented as:
 - Client can read chunks of various files using read command
 - Client provides the filename to read from
 - Client provides the chunk index (from bytes offset input by user)
 - Master looks for the file

Read

- Read is implemented as:
 - Client can read chunks of various files using read command
 - Client provides the filename to read from
 - Client provides the chunk index (from bytes offset input by user)
 - Master looks for the file
 - If found directs the address of closest chunk-server to the client

Read

- Read is implemented as:
 - Client can read chunks of various files using read command
 - Client provides the filename to read from
 - Client provides the chunk index (from bytes offset input by user)
 - Master looks for the file
 - If found directs the address of closest chunk-server to the client
 - Client communicated with the nearest chunk-server and receives the required data

Delete File

- Delete is implemented as:
 - Client can issue delete command

Delete File

- Delete is implemented as:
 - Client can issue delete command
 - Client provides the filename to delete

Delete File

- Delete is implemented as:
 - Client can issue delete command
 - Client provides the filename to delete
 - Master looks for the file and retrieves the chunk handle

Delete File

- Delete is implemented as:
 - Client can issue delete command
 - Client provides the filename to delete
 - Master looks for the file and retrieves the chunk handle
 - Master then removes it from the directory structure

Delete File

- Delete is implemented as:
 - Client can issue delete command
 - Client provides the filename to delete
 - Master looks for the file and retrieves the chunk handle
 - Master then removes it from the directory structure
 - The garbage cleaning mechanism then removes the file in various chunk-servers (which are now treated as orphaned chunks) on subsequent heartbeat (synchronization) messages

Append

- Append is implemented as:
 - Client can issue append command

Append

- Append is implemented as:
 - Client can issue append command
 - Master retrieves the chunk handle of the last chunk

Append

- Append is implemented as:
 - Client can issue append command
 - Master retrieves the chunk handle of the last chunk
 - The chunk servers containing this handle are retrieved and sent to client

Append

- Append is implemented as:
 - Client can issue append command
 - Master retrieves the chunk handle of the last chunk
 - The chunk servers containing this handle are retrieved and sent to client
 - All append information is pushed on these chunk-servers by the client

Append

- Append is implemented as:
 - Client can issue append command
 - Master retrieves the chunk handle of the last chunk
 - The chunk servers containing this handle are retrieved and sent to client
 - All append information is pushed on these chunk-servers by the client
 - Client issues write command which writes the pushed information on the chunk-servers

Snapshot

- Snapshot is implemented as:
 - Client can issue snapshot command

Snapshot

- Snapshot is implemented as:
 - Client can issue snapshot command
 - Client provides the directory name to snapshot

Snapshot

- Snapshot is implemented as:
 - Client can issue snapshot command
 - Client provides the directory name to snapshot
 - Master notes the timestamp and looks for the file in the directory

Snapshot

- Snapshot is implemented as:
 - Client can issue snapshot command
 - Client provides the directory name to snapshot
 - Master notes the timestamp and looks for the file in the directory
 - Master then checks for the nearest slave

Snapshot

- Snapshot is implemented as:
 - Client can issue snapshot command
 - Client provides the directory name to snapshot
 - Master notes the timestamp and looks for the file in the directory
 - Master then checks for the nearest slave
 - Master then dumps the directory structure and state information in a file

Snapshot

- Snapshot is implemented as:
 - Client can issue snapshot command
 - Client provides the directory name to snapshot
 - Master notes the timestamp and looks for the file in the directory
 - Master then checks for the nearest slave
 - Master then dumps the directory structure and state information in a file
 - This file is then dumped (written to) the chosen chunk-server

Chunk Replication

- Chunk-replication is done automatically when a chunk-server goes down or a new chunk-server is added

Chunk Replication

- Chunk-replication is done automatically when a chunk-server goes down or a new chunk-server is added
- It is implemented as:
 - Master retrieves the state information of the chunk-servers through the heartbeat messages

Chunk Replication

- Chunk-replication is done automatically when a chunk-server goes down or a new chunk-server is added
- It is implemented as:
 - Master retrieves the state information of the chunk-servers through the heartbeat messages
 - Master then decides the balancing strategy based on the number of chunks and the number of chunk-servers

Chunk Replication

- Chunk-replication is done automatically when a chunk-server goes down or a new chunk-server is added
- It is implemented as:
 - Master retrieves the state information of the chunk-servers through the heartbeat messages
 - Master then decides the balancing strategy based on the number of chunks and the number of chunk-servers
 - The chunks are then transferred between the chunk-servers based on the balancing strategy

Garbage Collection

- Garbage collection is done automatically when a chunk becomes orphaned

Garbage Collection

- Garbage collection is done automatically when a chunk becomes orphaned
- An orphaned chunk is a chunk whose entry is not in the master (because it was deleted)

Garbage Collection

- Garbage collection is done automatically when a chunk becomes orphaned
- An orphaned chunk is a chunk whose entry is not in the master (because it was deleted)
- It is implemented as:
 - Master retrieves the state information of the chunk-servers through the heartbeat messages

Garbage Collection

- Garbage collection is done automatically when a chunk becomes orphaned
- An orphaned chunk is a chunk whose entry is not in the master (because it was deleted)
- It is implemented as:
 - Master retrieves the state information of the chunk-servers through the heartbeat messages
 - Master then decides finds out the orphaned chunks based on this information

Garbage Collection

- Garbage collection is done automatically when a chunk becomes orphaned
- An orphaned chunk is a chunk whose entry is not in the master (because it was deleted)
- It is implemented as:
 - Master retrieves the state information of the chunk-servers through the heartbeat messages
 - Master then decides finds out the orphaned chunks based on this information
 - Master then issues a command to the chunk-server to delete the orphaned chunk in the respective chunk-server

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block
- It is implemented as:
 - When a chunk-server boots, it calculates the checksum of each chunk and stores it

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block
- It is implemented as:
 - When a chunk-server boots, it calculates the checksum of each chunk and stores it
 - Anytime there is a request to retrieve a particular chunk, the checksum for the particular chunk is calculated again

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block
- It is implemented as:
 - When a chunk-server boots, it calculates the checksum of each chunk and stores it
 - Anytime there is a request to retrieve a particular chunk, the checksum for the particular chunk is calculated again
 - In case this checksum does not match with the checksum initially calculate, master is informed

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block
- It is implemented as:
 - When a chunk-server boots, it calculates the checksum of each chunk and stores it
 - Anytime there is a request to retrieve a particular chunk, the checksum for the particular chunk is calculated again
 - In case this checksum does not match with the checksum initially calculate, master is informed
 - Master then retrieves the other chunk-servers that contain the same chunk

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block
- It is implemented as:
 - When a chunk-server boots, it calculates the checksum of each chunk and stores it
 - Anytime there is a request to retrieve a particular chunk, the checksum for the particular chunk is calculated again
 - In case this checksum does not match with the checksum initially calculate, master is informed
 - Master then retrieves the other chunk-servers that contain the same chunk
 - These chunk-servers then transfer the particular chunk to the primary chunk-server

Maintaining Integrity

- Integrity of various chunks in the system is maintained by using checksum
- Chunks are logically divided into 64KB block
- It is implemented as:
 - When a chunk-server boots, it calculates the checksum of each chunk and stores it
 - Anytime there is a request to retrieve a particular chunk, the checksum for the particular chunk is calculated again
 - In case this checksum does not match with the checksum initially calculate, master is informed
 - Master then retrieves the other chunk-servers that contain the same chunk
 - These chunk-servers then transfer the particular chunk to the primary chunk-server
 - The chunk-server then transfers the data to the client

Operation Log

- Operation log is maintained by the Master

Operation Log

- Operation log is maintained by the Master
- The following information is dumped to the log file

Operation Log

- Operation log is maintained by the Master
- The following information is dumped to the log file
 - File Name-space: Contains the directory structure

Operation Log

- Operation log is maintained by the Master
- The following information is dumped to the log file
 - File Name-space: Contains the directory structure
 - Meta-data: Contains the File-Hashname and the corresponding chunk handle and index

Operation Log

- Operation log is maintained by the Master
- The following information is dumped to the log file
 - File Name-space: Contains the directory structure
 - Meta-data: Contains the File-Hashname and the corresponding chunk handle and index
 - chunk db: list of all chunks (in all slaves)

Operation Log

- Operation log is maintained by the Master
- The following information is dumped to the log file
 - File Name-space: Contains the directory structure
 - Meta-data: Contains the File-Hashname and the corresponding chunk handle and index
 - chunk db: list of all chunks (in all slaves)
 - slave list: list of active slaves

Operation Log

- Operation log is maintained by the Master
- The following information is dumped to the log file
 - File Name-space: Contains the directory structure
 - Meta-data: Contains the File-Hashname and the corresponding chunk handle and index
 - chunk db: list of all chunks (in all slaves)
 - slave list: list of active slaves
 - snapshot record: timestamp, IP and port information from where snapshot was invoked

Conclusion

Conclusion

- Google File System can be implemented to support large files through chunk mechanism

Conclusion

- Google File System can be implemented to support large files through chunk mechanism
- The master serves as the main control node and communicator between client and chunks

Conclusion

- Google File System can be implemented to support large files through chunk mechanism
- The master serves as the main control node and communicator between client and chunks
- Client can issue various commands to perform different operations on files

Conclusion

- Google File System can be implemented to support large files through chunk mechanism
- The master serves as the main control node and communicator between client and chunks
- Client can issue various commands to perform different operations on files
- The heartbeat messages provide an efficient synchronization mechanism between the master and the chunk-servers

Conclusion

- Google File System can be implemented to support large files through chunk mechanism
- The master serves as the main control node and communicator between client and chunks
- Client can issue various commands to perform different operations on files
- The heartbeat messages provide an efficient synchronization mechanism between the master and the chunk-servers
- The garbage cleaning mechanism ensures that no redundant chunks are present in any chunk-server

Conclusion

- Google File System can be implemented to support large files through chunk mechanism
- The master serves as the main control node and communicator between client and chunks
- Client can issue various commands to perform different operations on files
- The heartbeat messages provide an efficient synchronization mechanism between the master and the chunk-servers
- The garbage cleaning mechanism ensures that no redundant chunks are present in any chunk-server
- The load-balancing mechanism ensures equitable distribution of the chunks stored by the various chunk-servers

Future Scope

Future Scope

- Some recent changes[2] made to GFS by google include:

Future Scope

- Some recent changes[2] made to GFS by google include:
 - New applications now require low latency

Future Scope

- Some recent changes[2] made to GFS by google include:
 - New applications now require low latency
 - Multiple GFS masters now on a pool of chunk-servers

Future Scope

- Some recent changes[2] made to GFS by google include:
 - New applications now require low latency
 - Multiple GFS masters now on a pool of chunk-servers
 - Multi-cell approach with master for each cell

Future Scope

- Some recent changes[2] made to GFS by google include:
 - New applications now require low latency
 - Multiple GFS masters now on a pool of chunk-servers
 - Multi-cell approach with master for each cell
 - More than one cell per data center

Future Scope

- Some recent changes[2] made to GFS by google include:
 - New applications now require low latency
 - Multiple GFS masters now on a pool of chunk-servers
 - Multi-cell approach with master for each cell
 - More than one cell per data center
 - Cells across network function as distinct file systems

Future Scope

- Some recent changes[2] made to GFS by google include:
 - New applications now require low latency
 - Multiple GFS masters now on a pool of chunk-servers
 - Multi-cell approach with master for each cell
 - More than one cell per data center
 - Cells across network function as distinct file systems
- Implementation can be extended to incorporate these changes

References

References

[1] The Google File System; Sanjay Ghemawat et al.

<http://www.cs.cornell.edu/courses/cs614/2004sp/papers/gfs.pdf>

[2] Google File System

<http://google-file-system.wikispaces.asu.edu/>

Thank You!